

Deep Learning - Introduction

Itthi Chatnuntawech

ARTIFICIAL INTELLIGENCE

Any technique that enables computers to mimic human behavior



1950's 1960's 1970's 1980's

1990's 2000's 2010's

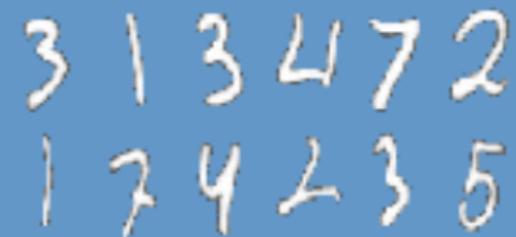
MACHINE LEARNING

Ability to learn without explicitly being programmed

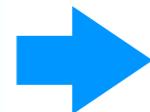


DEEP LEARNING

Learn underlying features in data using neural networks



1950's 1960's 1970's 1980's 1990's 2000's 2010's



Fasting blood sugar test



x mg/dL

normal if $x < 100$

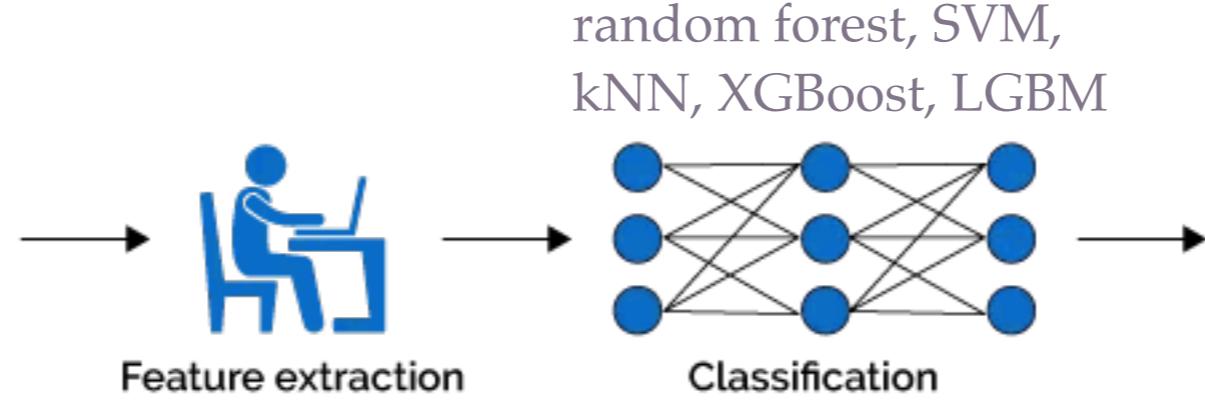
prediabetes if $100 \leq x \leq 125$

diabetes if $125 < x$

Example: Watermelon or Orange?

Traditional Machine Learning

Input image

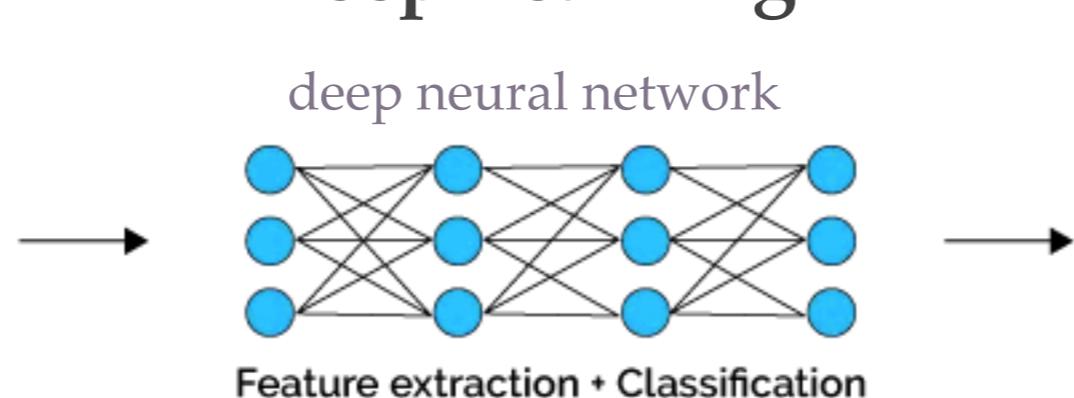


Output

watermelon

Deep Learning

Input image



Output

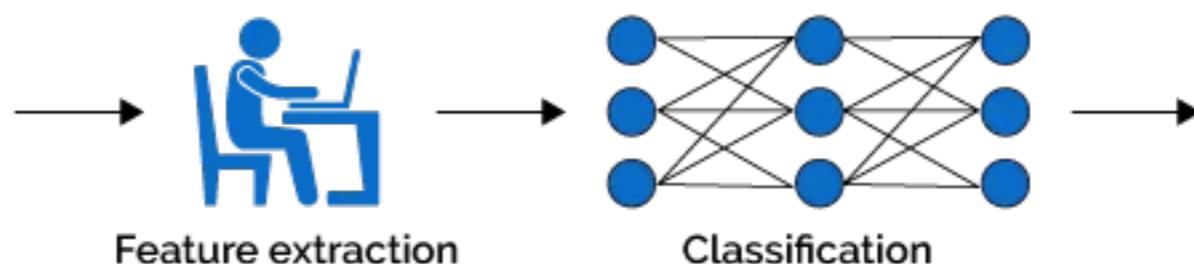
watermelon

Some parts taken from: Azua Tech

Watermelon or Orange?

Traditional Machine Learning

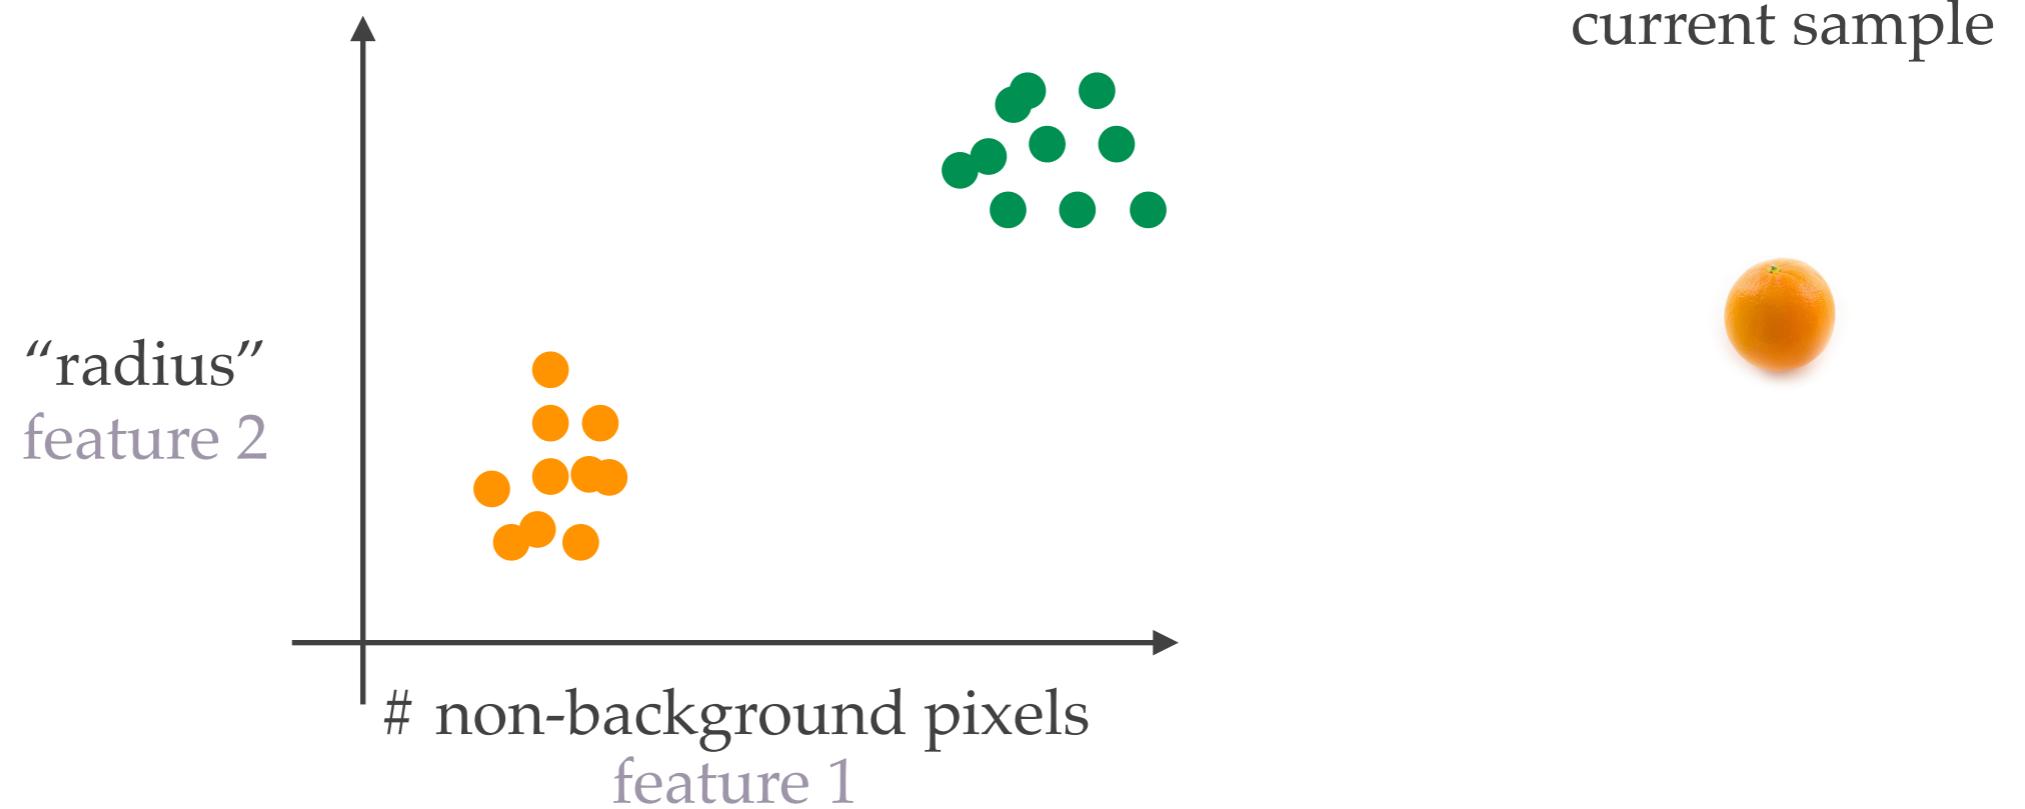
Input image



Output

watermelon

Phase: Training

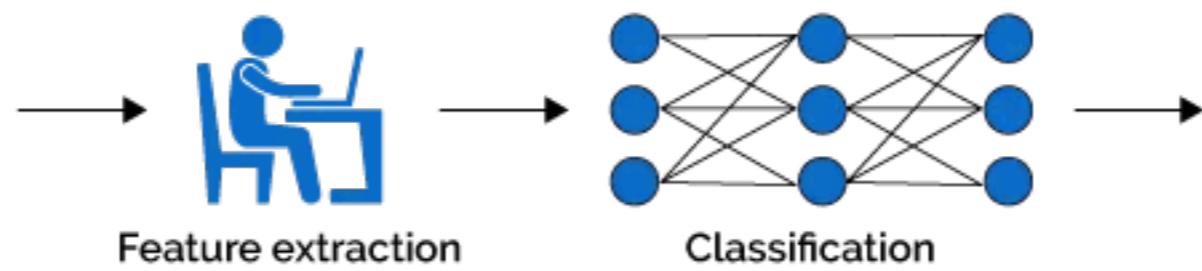


keywords: feature extraction, feature engineering, handcrafted / hand-designed features, feature space, descriptors

Watermelon or Orange?

Traditional Machine Learning

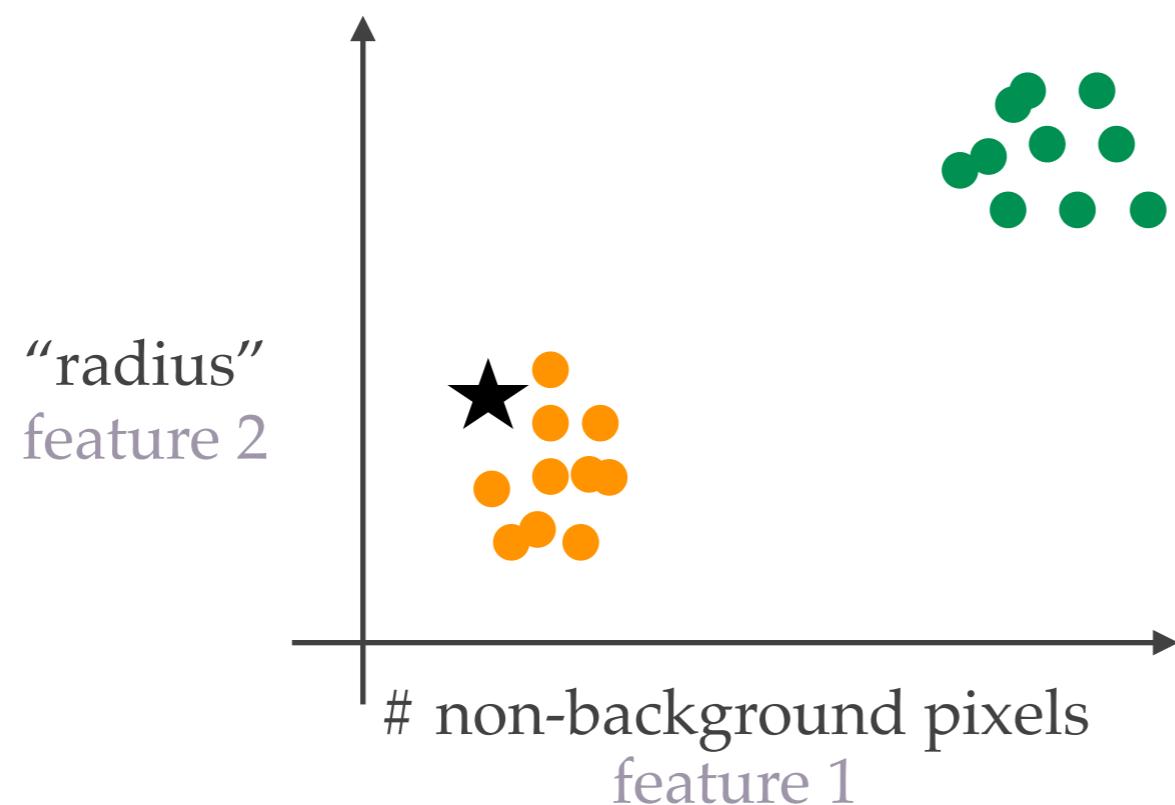
Input image



Output

watermelon

Phase: Testing



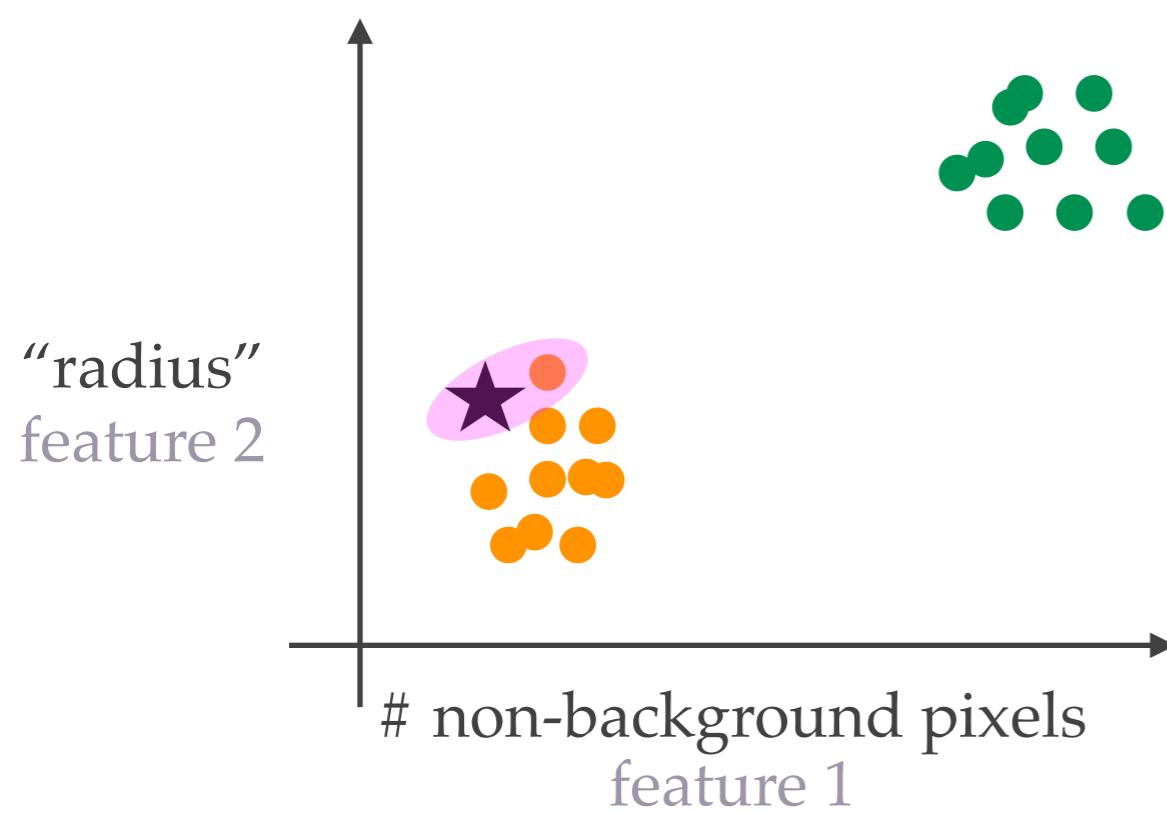
current sample



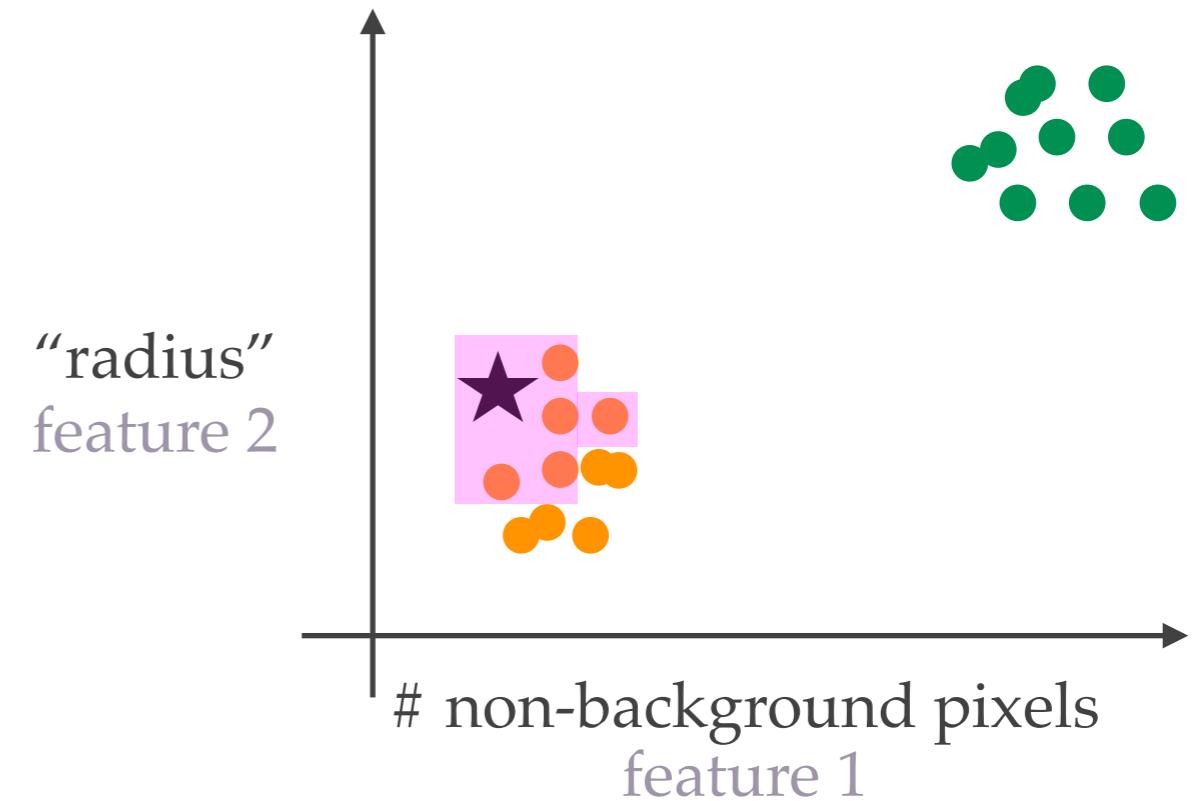
Orange!

keyword: decision boundary

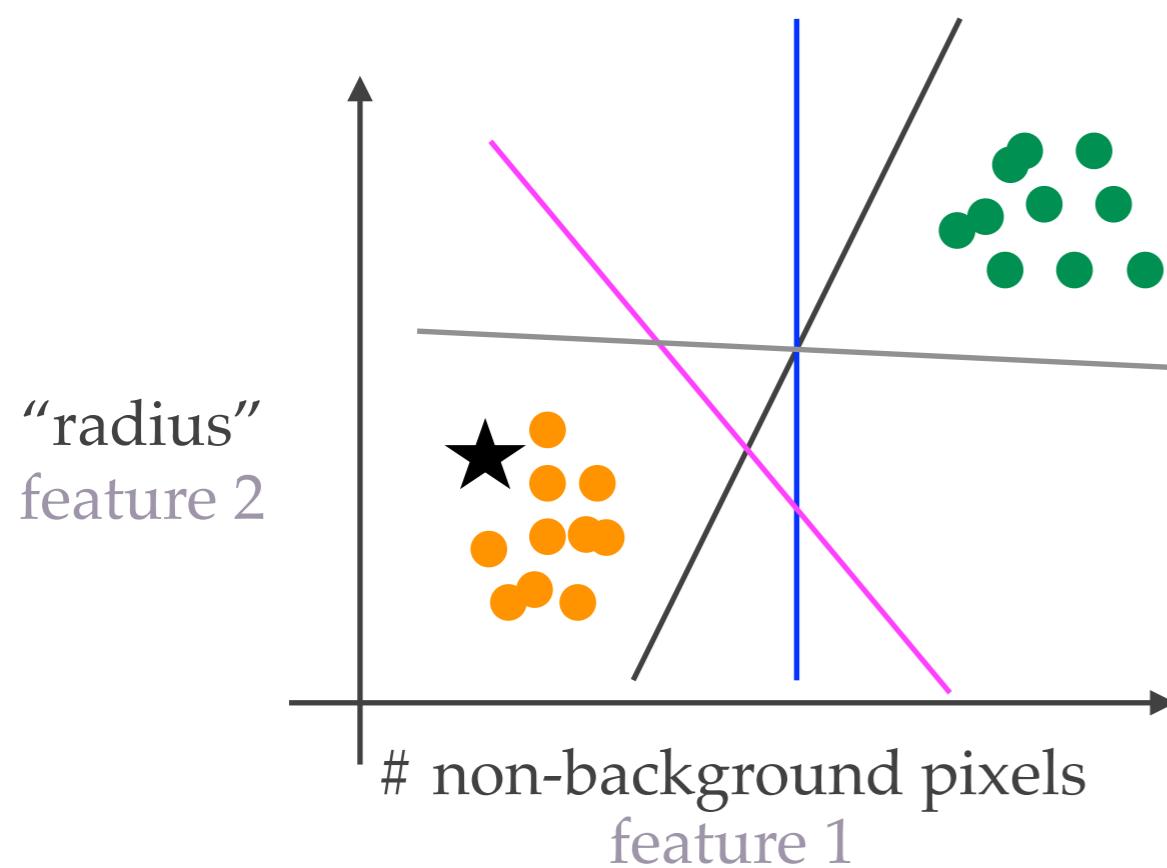
Nearest Neighbor



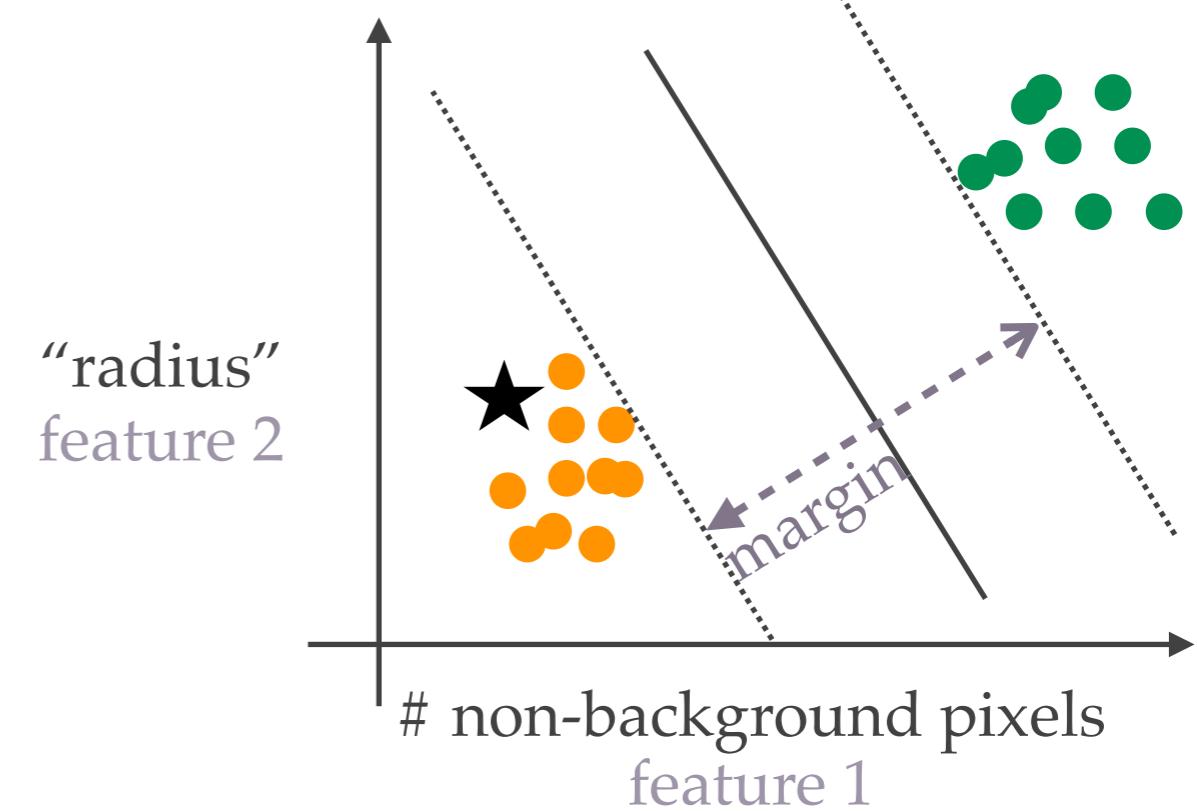
k-Nearest Neighbors (kNN)



Linear Classifier



Support Vector Machine (SVM) with a linear kernel



Watermelon or Orange?

Traditional Machine Learning

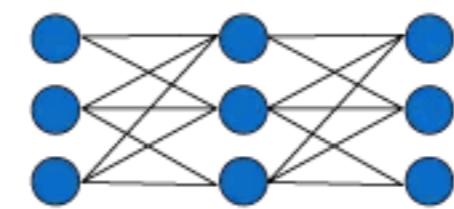
Input image



hand-engineered/
hand-crafted features



random forest, SVM,
kNN, XGBoost, LGBM



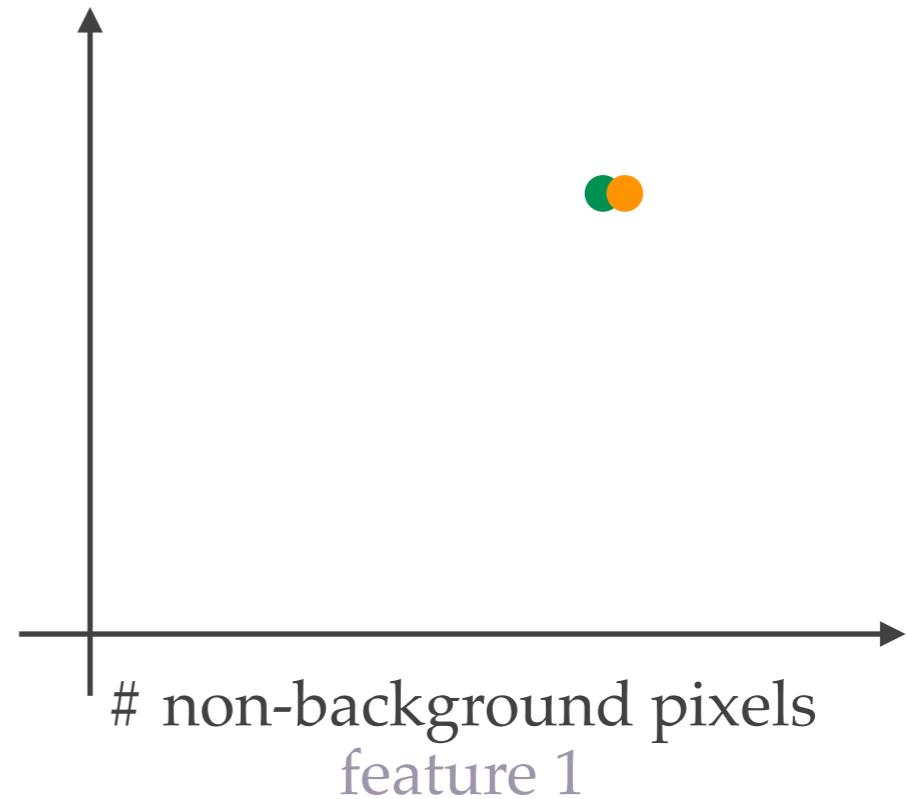
Output

watermelon

Problem with feature engineering?



“radius”
feature 2

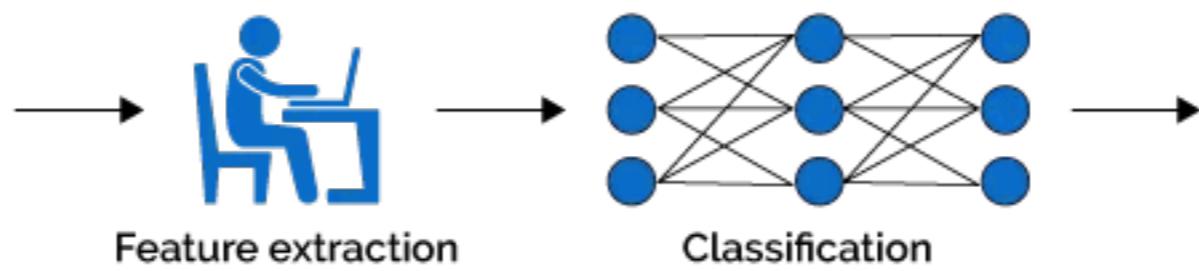


Some parts taken from: Azua Tech

Watermelon or Orange?

Traditional Machine Learning

Input image

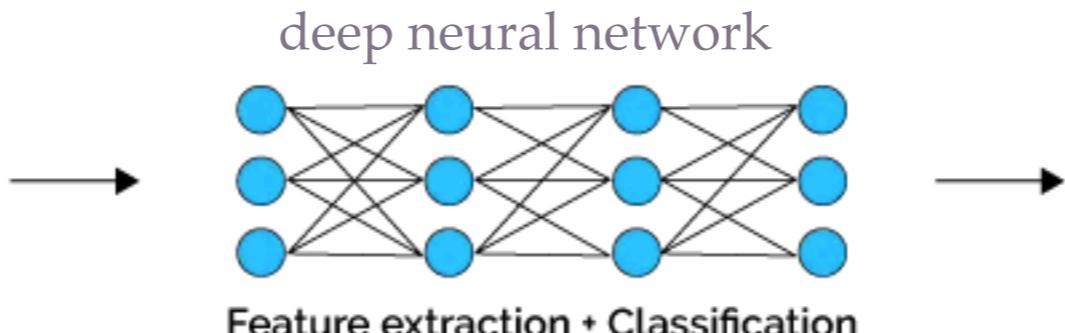


Output

watermelon

Deep Learning

Input image



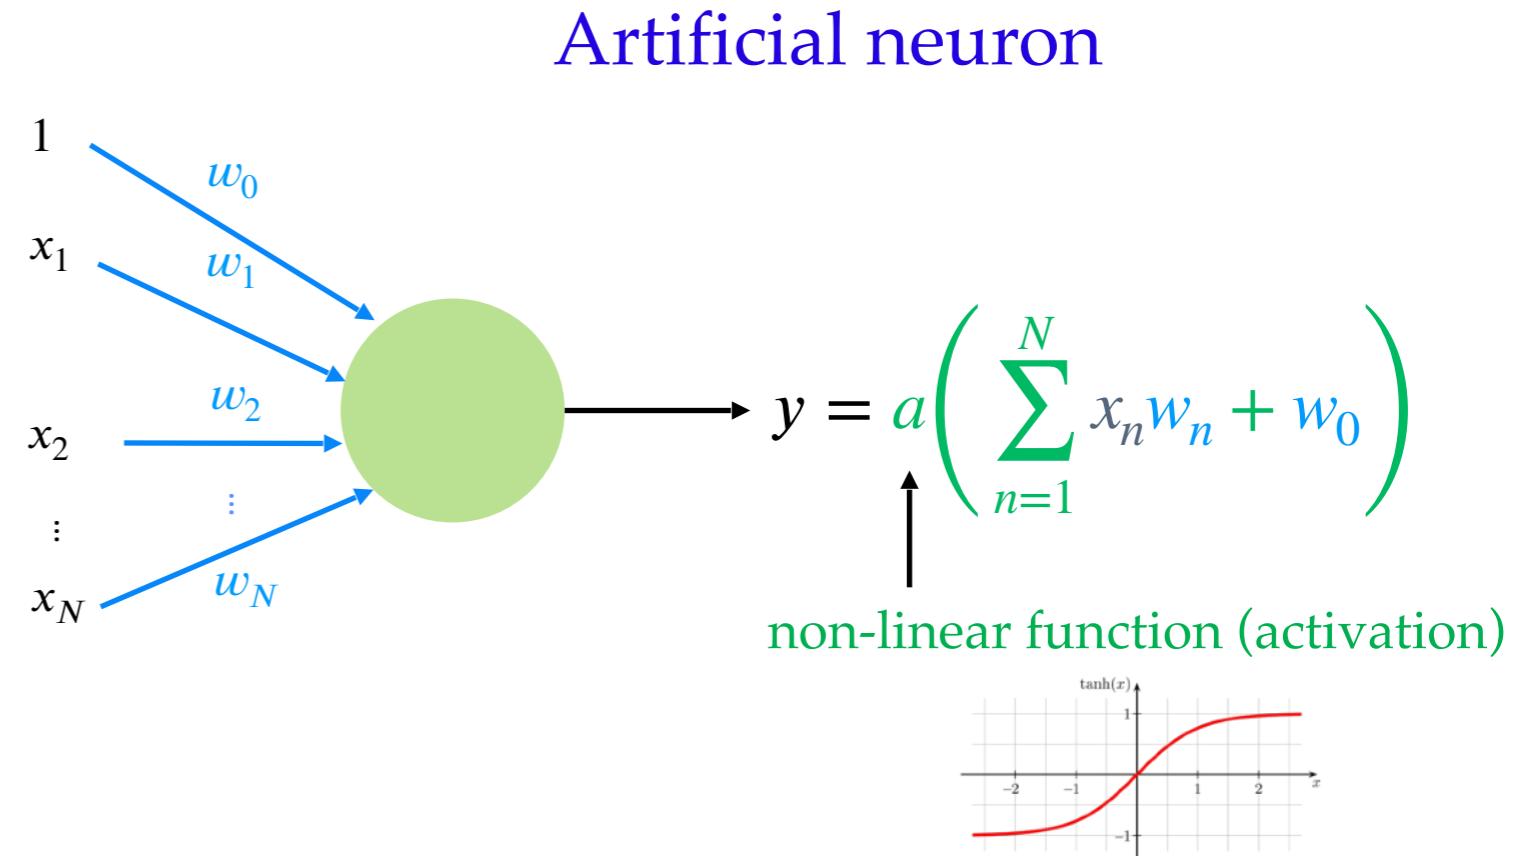
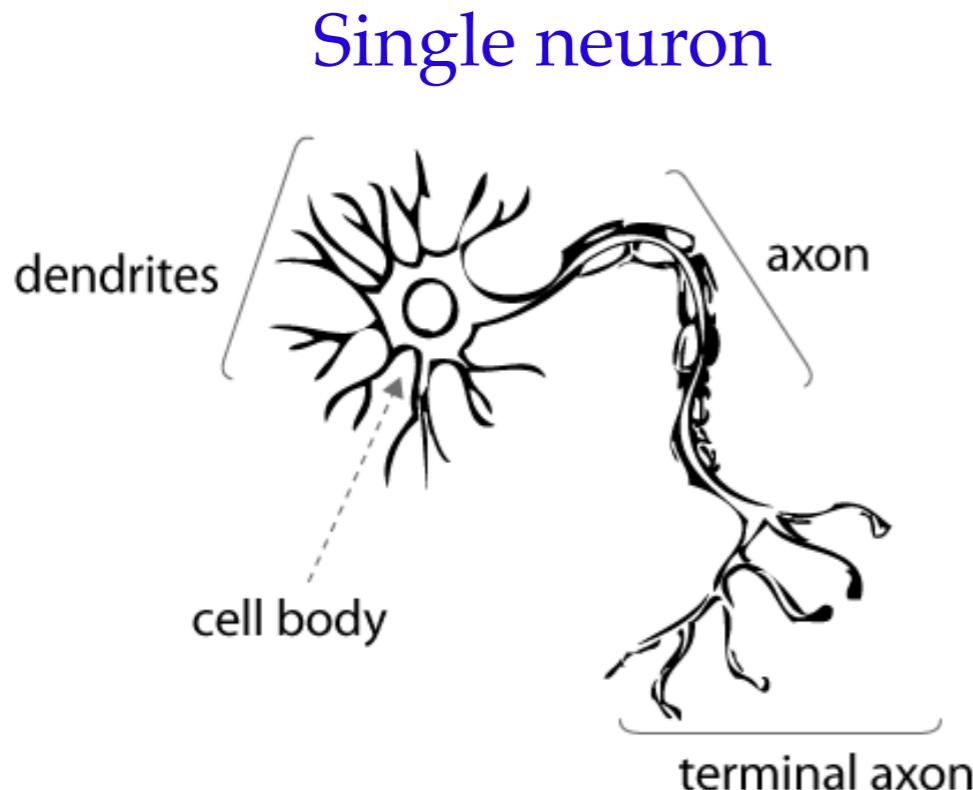
Output

watermelon

Some parts taken from: Azua Tech

Deep Learning

- ❖ Deep learning is a subfield of machine learning that stems from artificial neural networks (ANN)



If we set all the weights to 0, then $y = 0$.

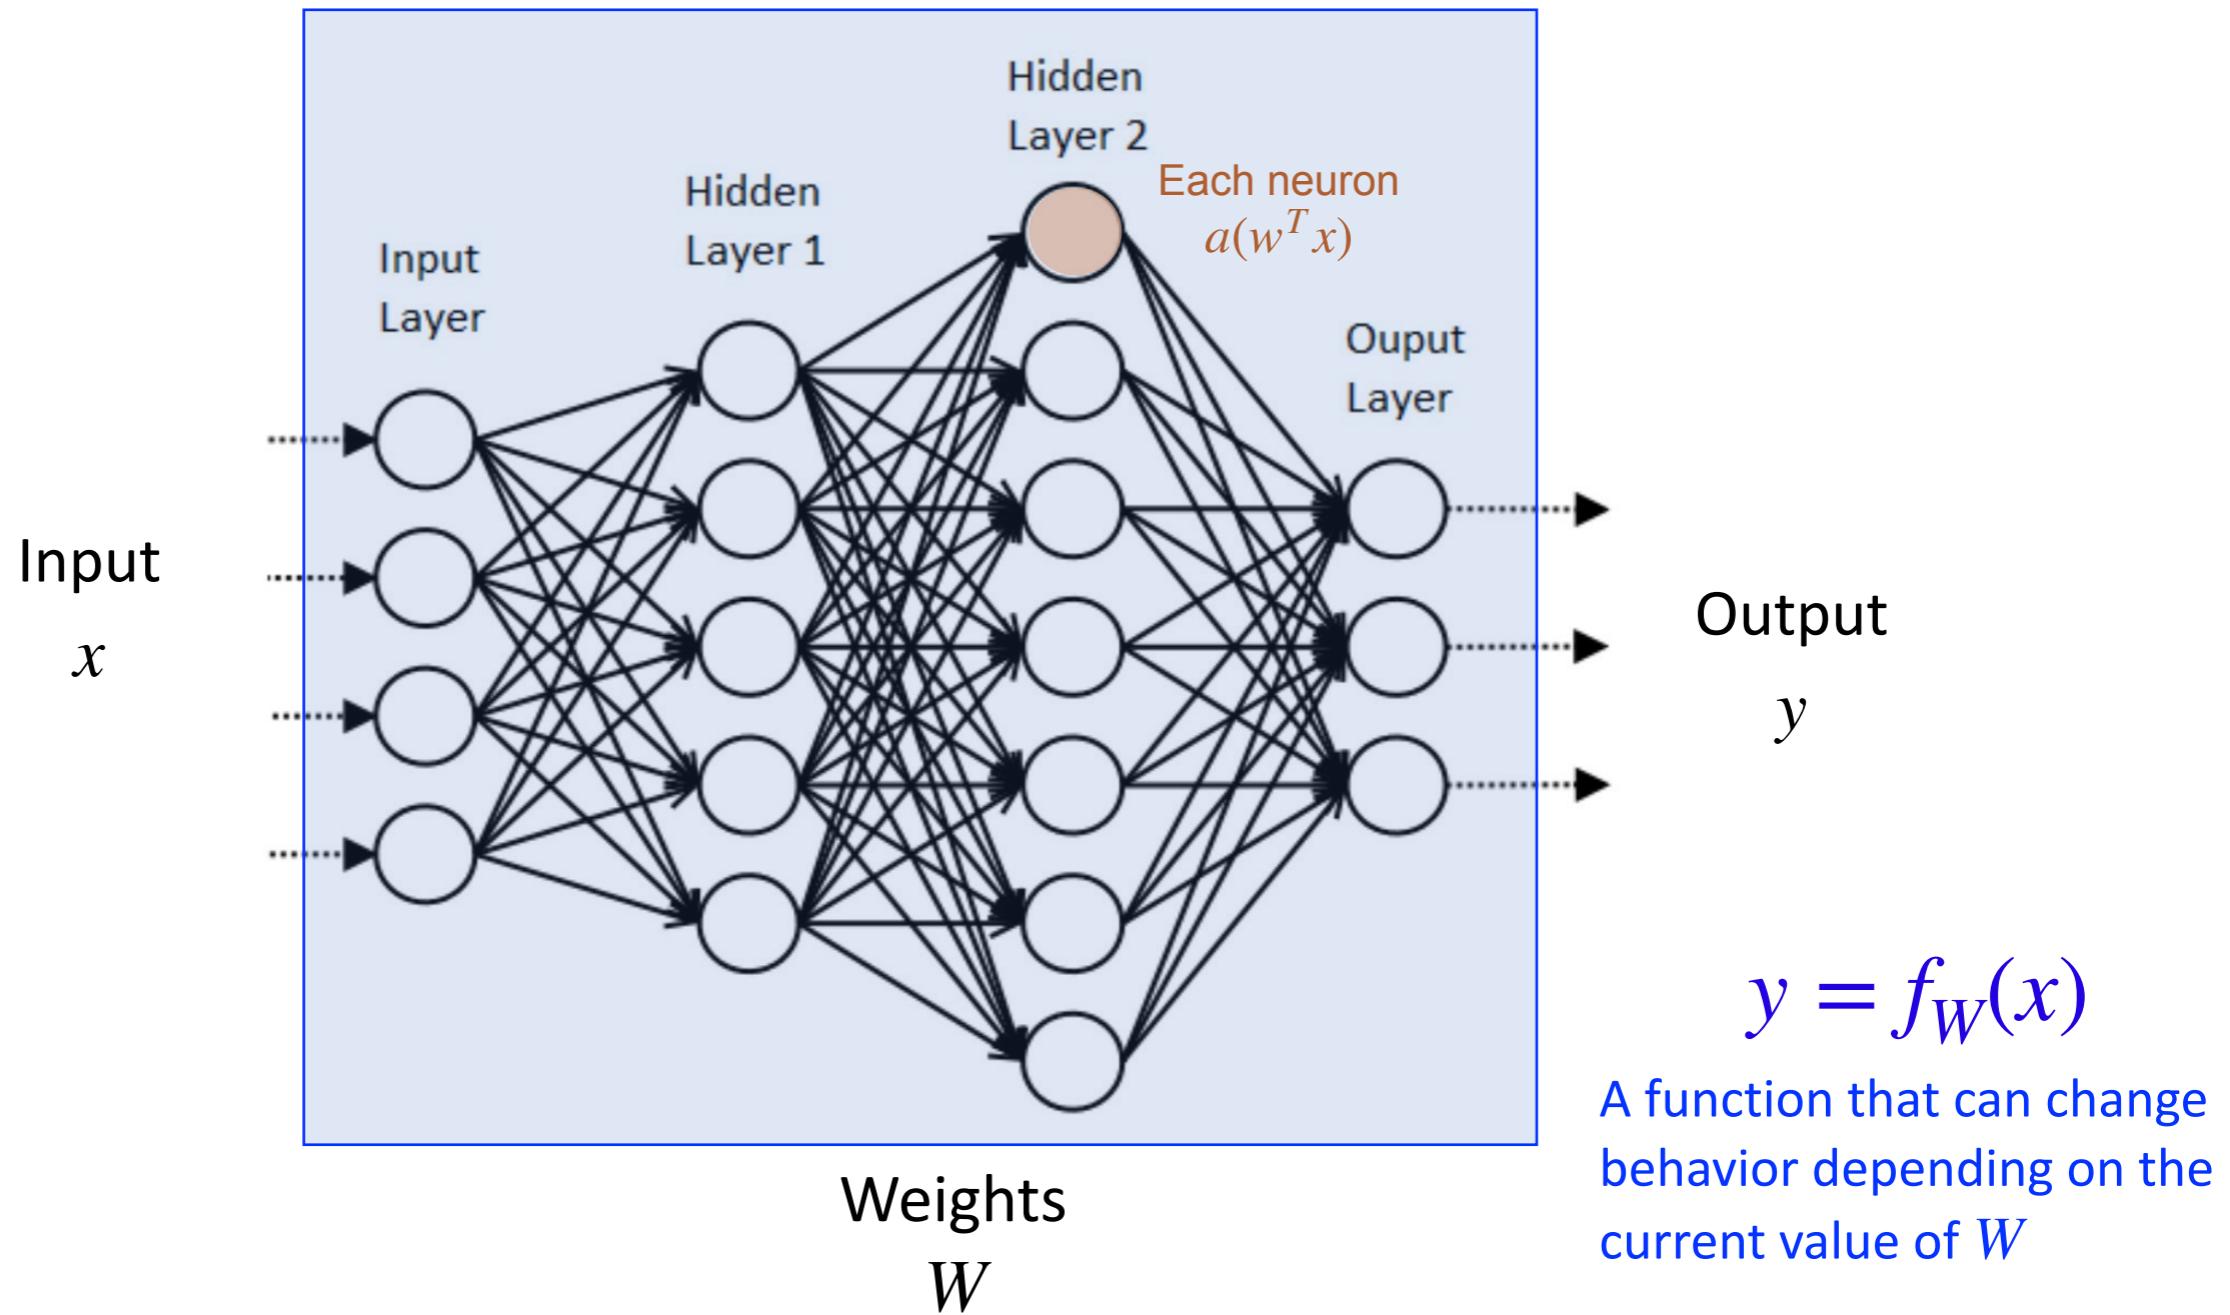
If we set $w_1 = 1$ and the rest to 0, then $y = a(x_1)$.

If we set $w_0 = 0$ and the rest to 1, then $y = a(x_1 + x_2 + \dots + x_N)$.

Different weights give rise
to different behavior

Deep Learning

- An artificial neural network with **many hidden layers** is called a **deep** artificial neural network (ANN)

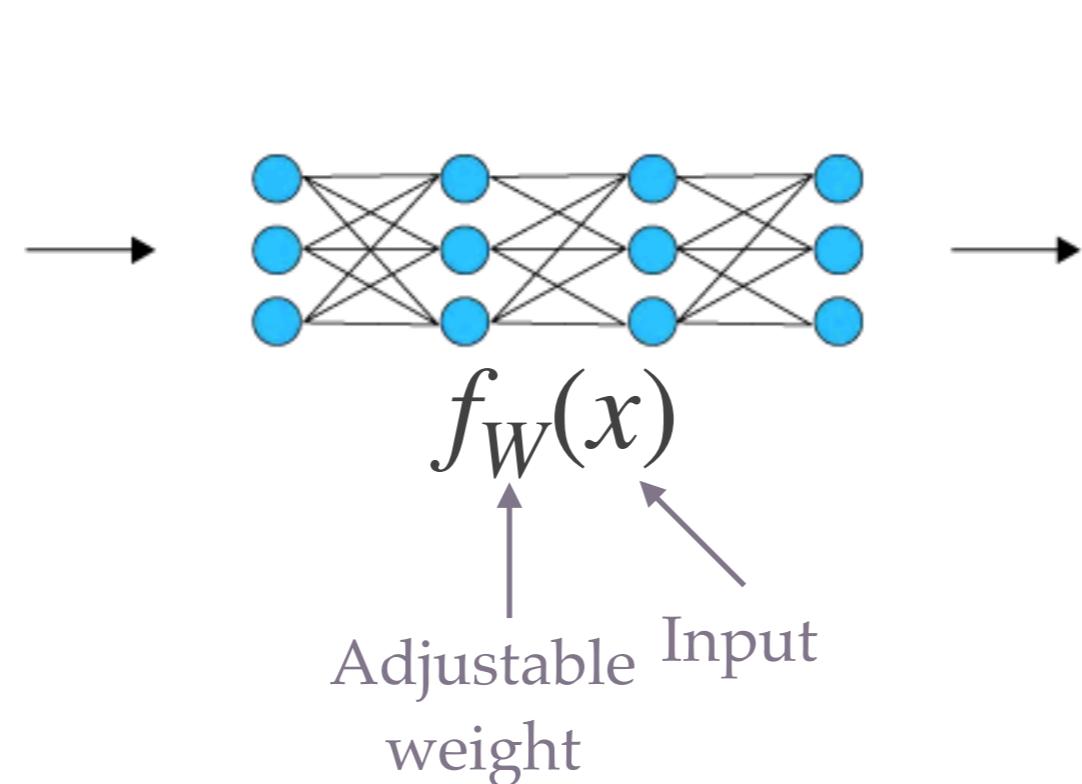


Deep Learning

- ❖ Think of a deep NN as a universal function approximator

Image classification

Input image

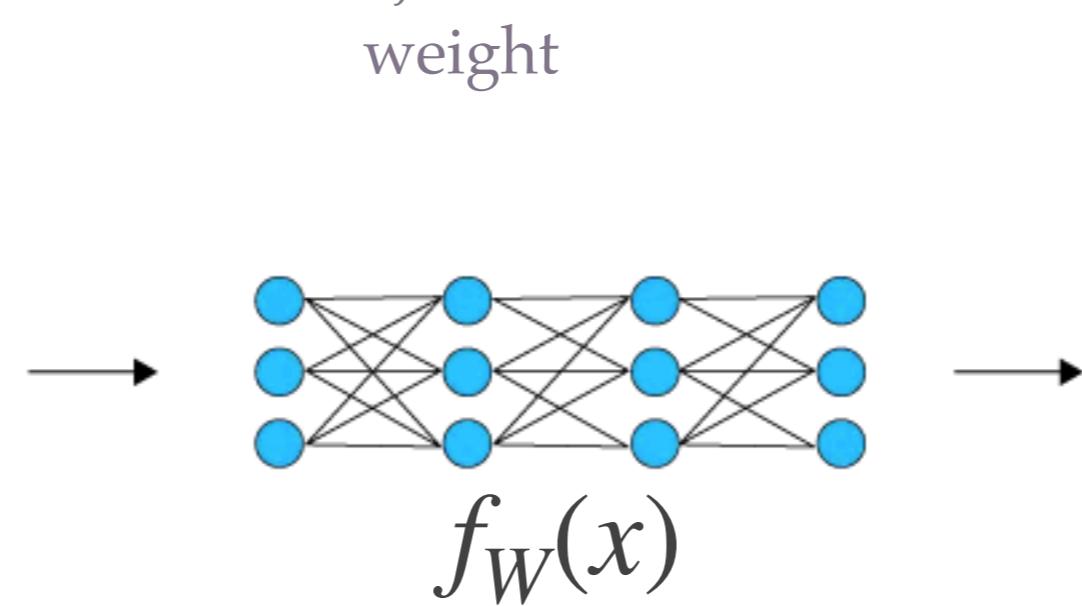


Output: Label

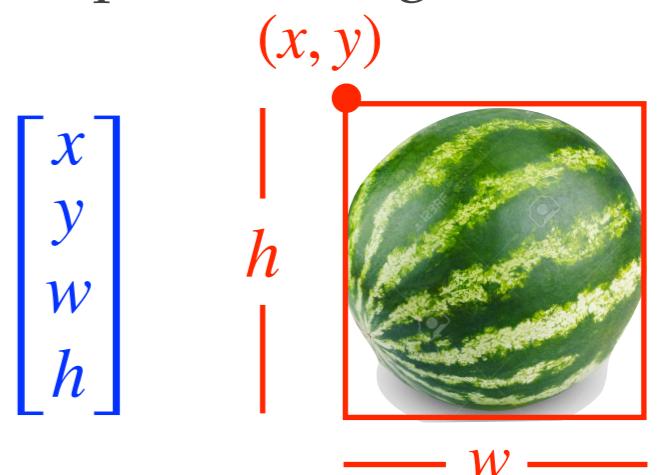
1 (watermelon)

Object detection

Input image



Output: rectangular box



Fully connected layer
(Dense)

Convolutional layer
Conv1D, 2D, 3D, ...
separable Conv

Optimizer
SGD
Adam
RMSprop

Evaluation metric
accuracy
F1-score
AUC
confusion matrix

Loss function
categorical crossentropy
binary crossentropy
mean squared error
mean absolute error

Regularization
Dropout
Data augmentation
 l_1, l_2 regularizations

Pooling layer
max-pooling
average-pooling

Activation function
sigmoid
softmax
ESP (swish)
ReLU



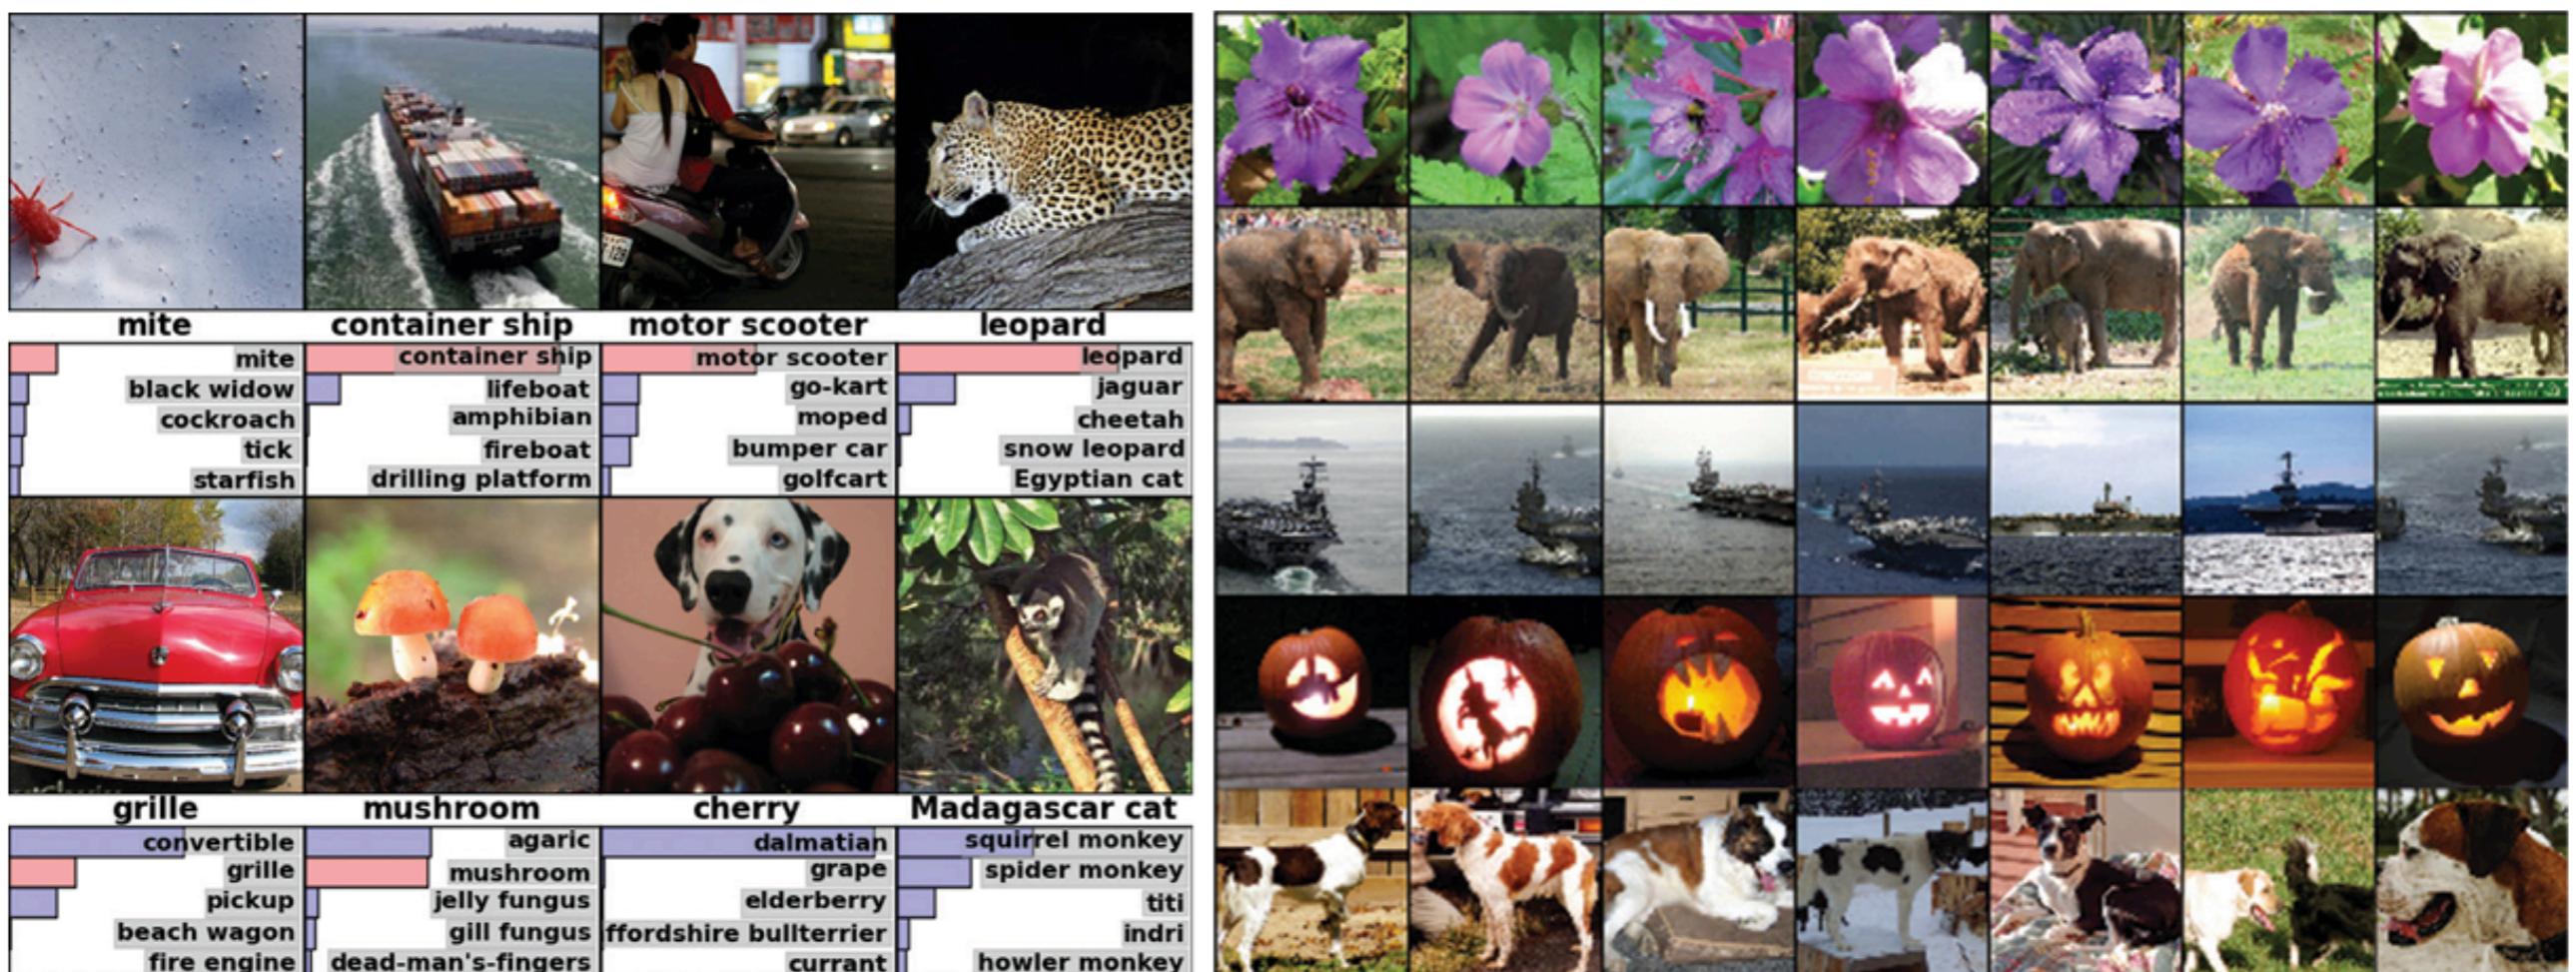
- ❖ **Combine basic components to build a neural network**
 - More components → “More” representative power

ImageNet

IMAGENET

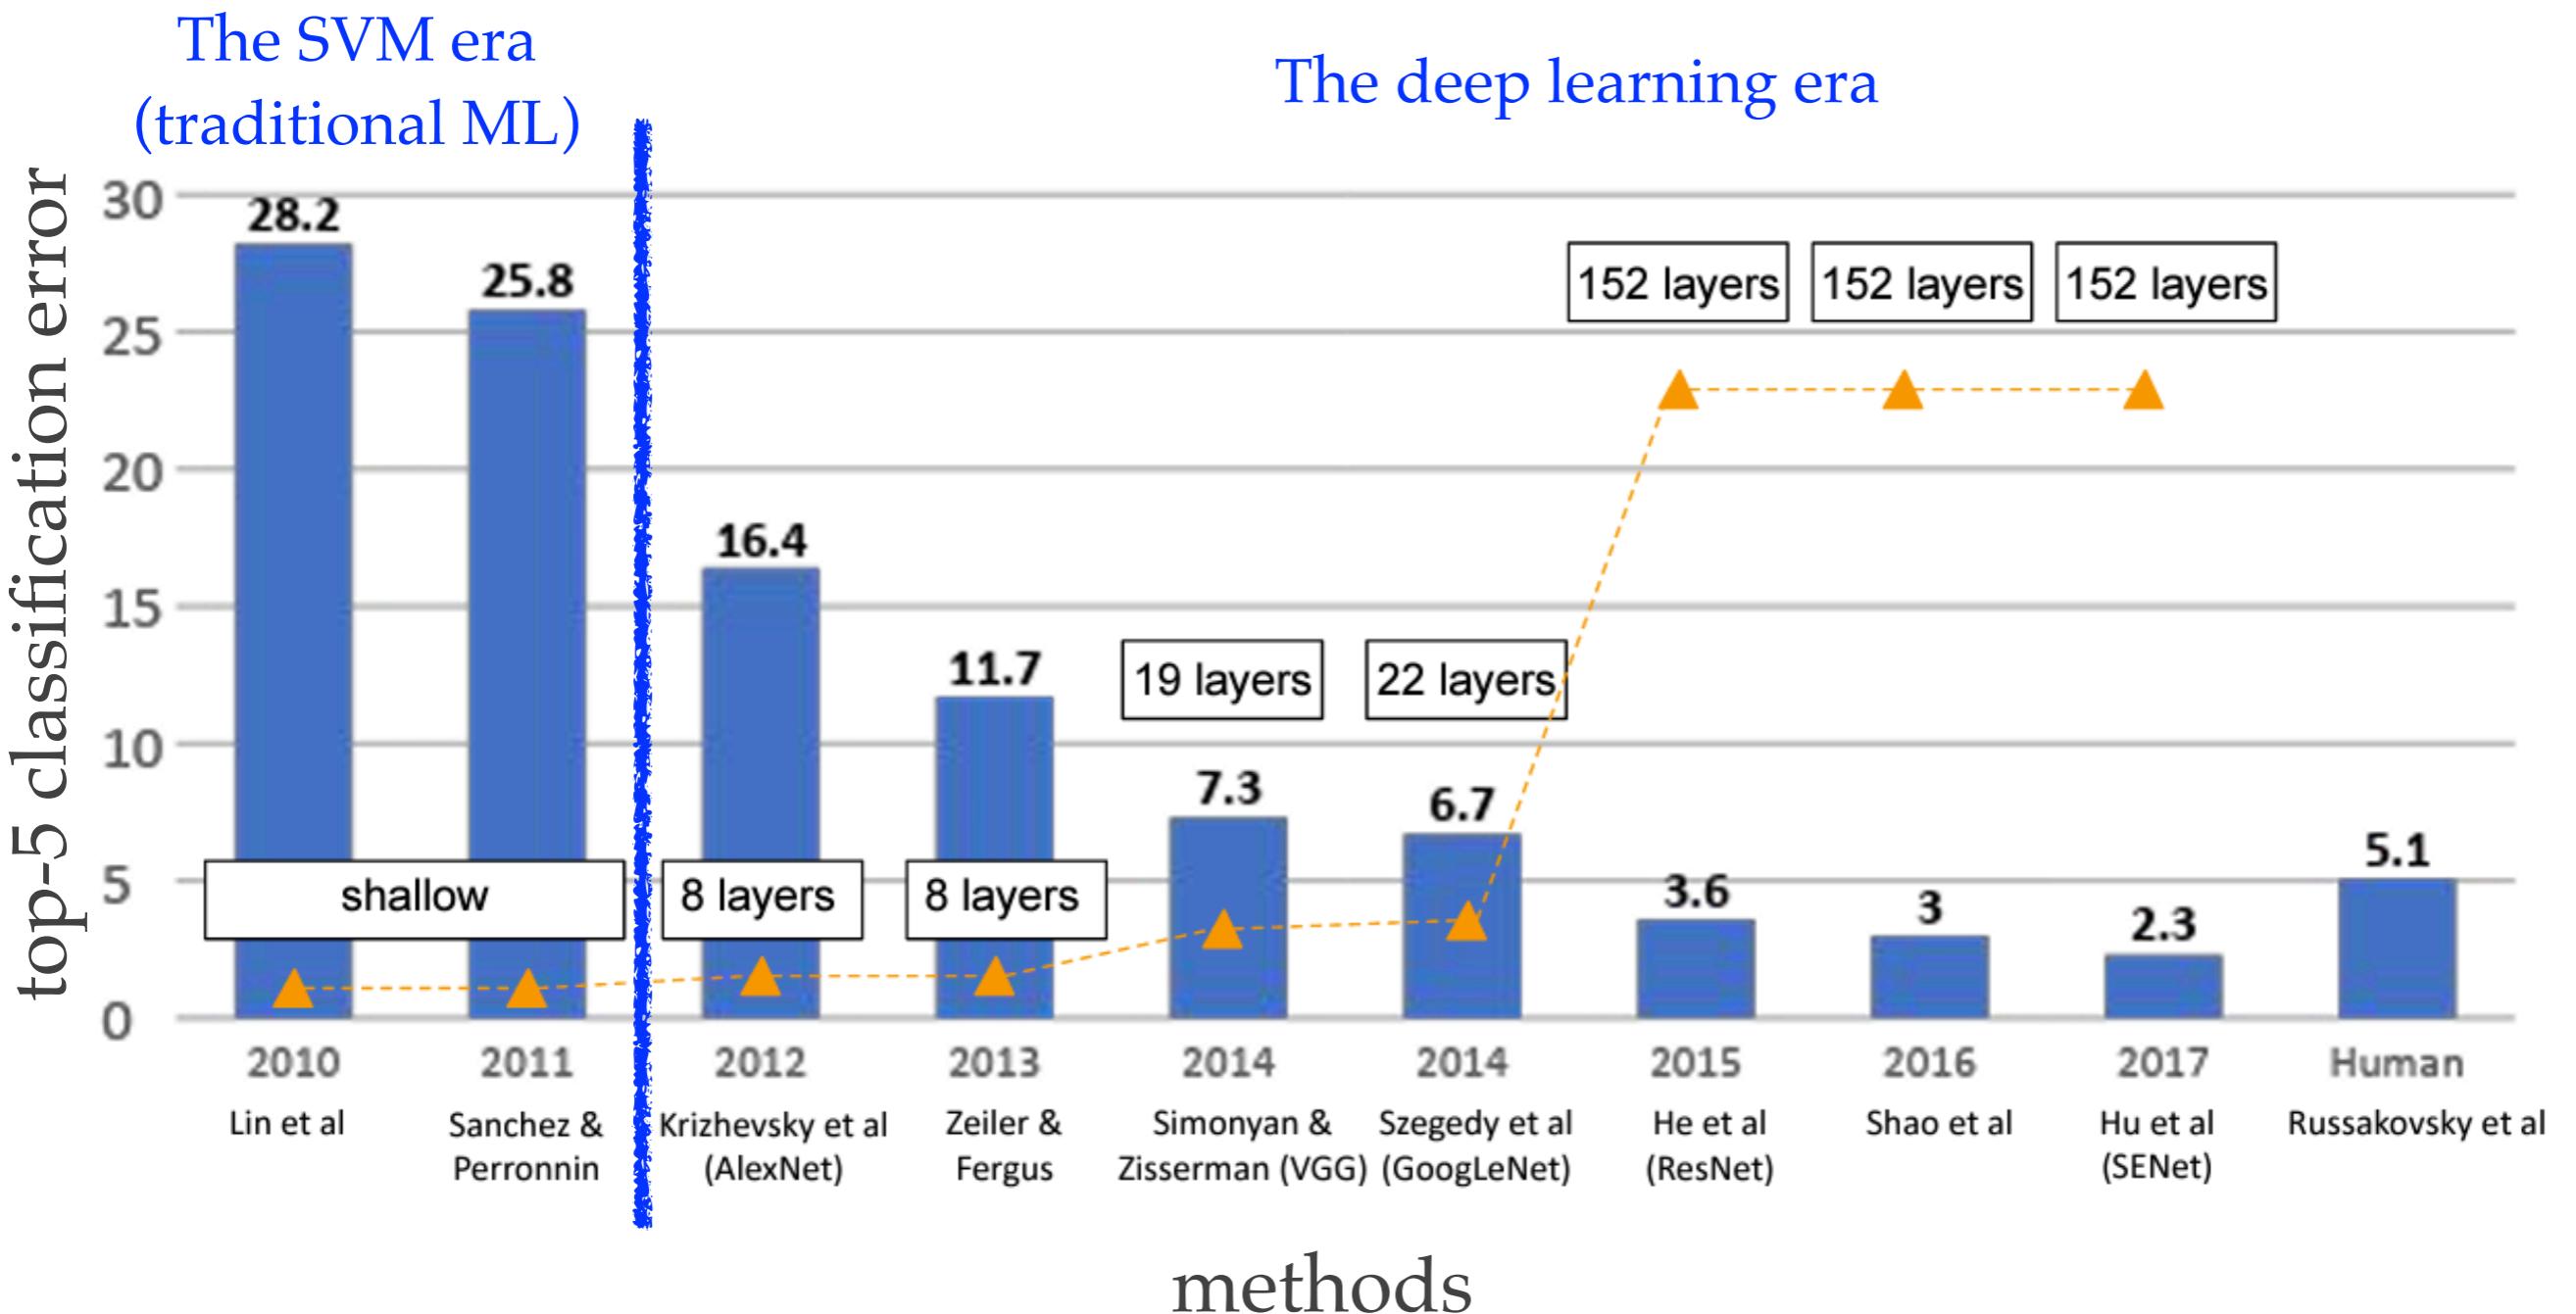
1000 classes, 14,197,122 images

Figure 4. (Left) Eight ILSVRC-2010 test images and the five labels considered most probable by our model. The correct label is written under each image, and the probability assigned to the correct label is also shown with a red bar (if it happens to be in the top 5). (Right) Five ILSVRC-2010 test images in the first column. The remaining columns show the six training images that produce feature vectors in the last hidden layer with the smallest Euclidean distance from the feature vector for the test image.



ImageNet

ImageNet Large Scale Visual Recognition Challenge (ILSVRC) winners



Supervised Model Training

Step 1: Create a neural network with some initial weight f_{W_0}

Neural network



Training iteration 0
 f_{W_0}

Supervised Model Training

Step 2: Prepare a dataset which is a collection of input-output pairs

Prepared inputs

x



Prepared outputs
(true labels)

y

1

Neural network



Training iteration 0

f_{W_0}



0



0

orange: 0
watermelon: 1

1

Supervised Model Training

Step 3: Pass the prepared inputs to the network

Prepared inputs

x



Estimated outputs

$$\hat{y} = f_{W_0}(x)$$

0.6

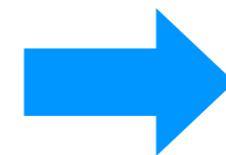
Prepared outputs
(true labels)

y

1



Neural network



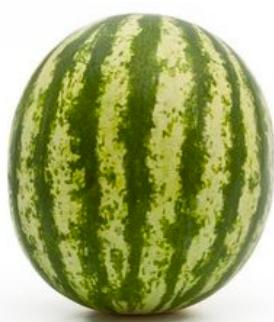
0.6



Training iteration 0
 f_{W_0}

0.5

0



0.5

orange: 0

watermelon: 1

1

Supervised Model Training

Step 4: Compare \hat{y} to y and modify the weights of the neural network to make \hat{y} approach y using the backpropagation algorithm

Prepared inputs



x

A simple loss function: $L(y, \hat{y}) = \sum_{i=1}^N (\hat{y}_i - y_i)^2$

$$L(y, \hat{y}) = (0.6 - 1)^2 + (0.6 - 0)^2 + (0.5 - 0)^2 + (0.5 - 1)^2 = 1.02$$

Estimated outputs
 $\hat{y} = f_{W_0}(x)$

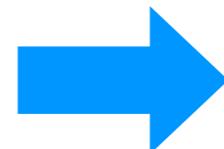
Prepared outputs
(true labels)

y

0.6

1

Neural network



Training iteration 0

f_{W_0}

Training loss

Iteration



0.6

0

0.5

0

0.5

orange: 0

watermelon: 1

Supervised Model Training

Repeat steps 3 and 4 to continuously improve the weights of the neural network

Prepared inputs

x



A simple loss function: $L(y, \hat{y}) = \sum_{i=1}^N (\hat{y}_i - y_i)^2$

$$L(y, \hat{y}) = (0.6 - 1)^2 + (0.4 - 0)^2 + (0.4 - 0)^2 + (0.6 - 1)^2 = 0.64$$

Estimated outputs
 $\hat{y} = f_{W_1}(x)$

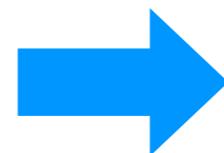
0.6

1

Prepared outputs
(true labels)

y

1



Neural network



0.4

0

0.4

0

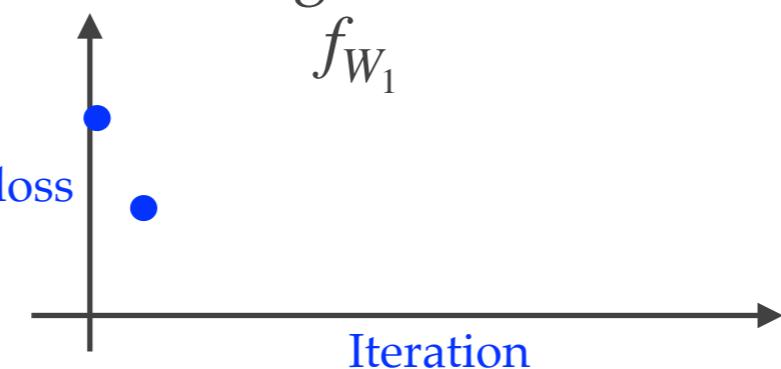


Training iteration 1

f_{W_1}



Training loss



0.6

1

orange: 0
watermelon: 1

Supervised Model Training

Repeat steps 3 and 4 to continuously improve the weights of the neural network

Prepared inputs

x



A simple loss function: $L(y, \hat{y}) = \sum_{i=1}^N (\hat{y}_i - y_i)^2$

$$L(y, \hat{y}) = (0.7 - 1)^2 + (0.35 - 0)^2 + (0.4 - 0)^2 + (0.6 - 1)^2 = 0.5325$$

Estimated outputs
 $\hat{y} = f_{W_2}(x)$

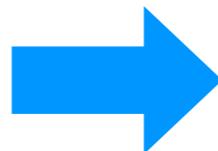
Prepared outputs
(true labels)

y

0.7

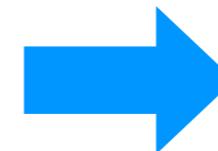
1

Neural network



0.35

0



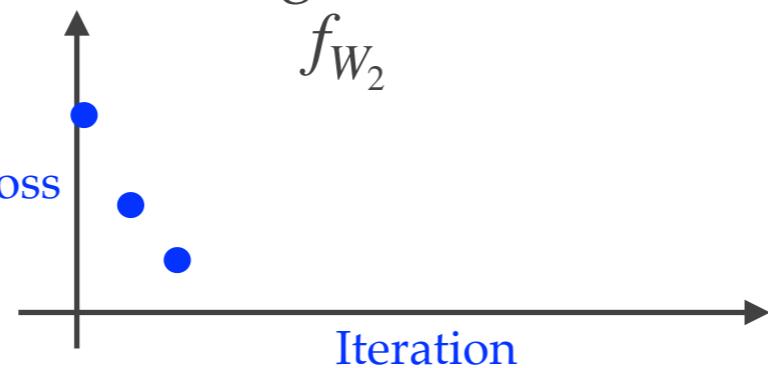
0.4

0

Training iteration 2

f_{W_2}

Training loss



0.6

1

orange: 0
watermelon: 1

Supervised Model Training

Repeat steps 3 and 4 to continuously improve the weights of the neural network

Prepared inputs
 x



A simple loss function: $L(y, \hat{y}) = \sum_{i=1}^N (\hat{y}_i - y_i)^2$

$$L(y, \hat{y}) = (0.8 - 1)^2 + (0.15 - 0)^2 + (0.1 - 0)^2 + (0.9 - 1)^2 = 0.0825$$

Estimated outputs
 $\hat{y} = f_{W_{50}}(x)$

0.8

1

Prepared outputs
(true labels)
 y

1



Neural network



0.15

0

0.1

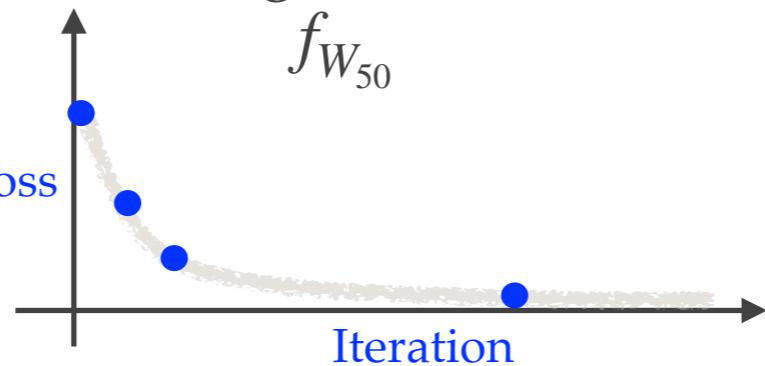
0



Training iteration 50

$f_{W_{50}}$

Training loss



0.9

1

orange: 0
watermelon: 1

Supervised Model Training

Repeat steps 3 and 4 to continuously improve the weights of the neural network

Prepared inputs
 x



A simple loss function: $L(y, \hat{y}) = \sum_{i=1}^N (\hat{y}_i - y_i)^2$

$$L(y, \hat{y}) = (0.99 - 1)^2 + (0.05 - 0)^2 + (0.01 - 0)^2 + (0.97 - 1)^2 = 0.0036$$

Estimated outputs
 $\hat{y} = f_{W_{1000}}(x)$

0.99

Prepared outputs
(true labels)
 y

1



Neural network



0.05



0

Training iteration 1000

$f_{W_{1000}}$



Training loss

Iteration

0.01

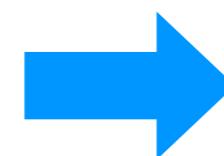
0

0.97

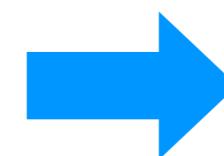
orange: 0
watermelon: 1

Test the Trained Model

Test image (unseen)



Trained neural network



Predicted Label

0.01

less than 0.5, so it is
likely orange

orange: 0
watermelon: 1

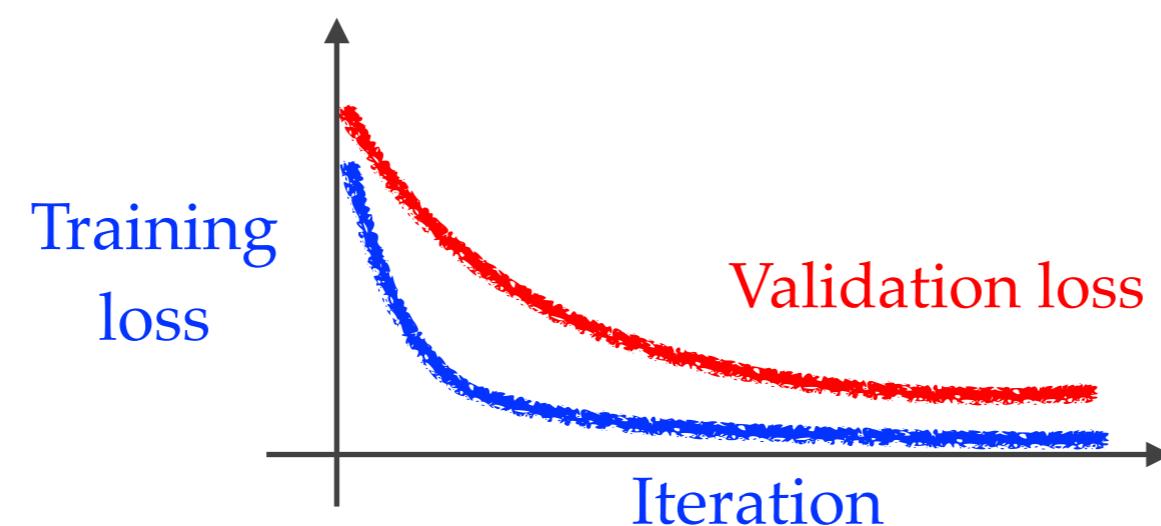
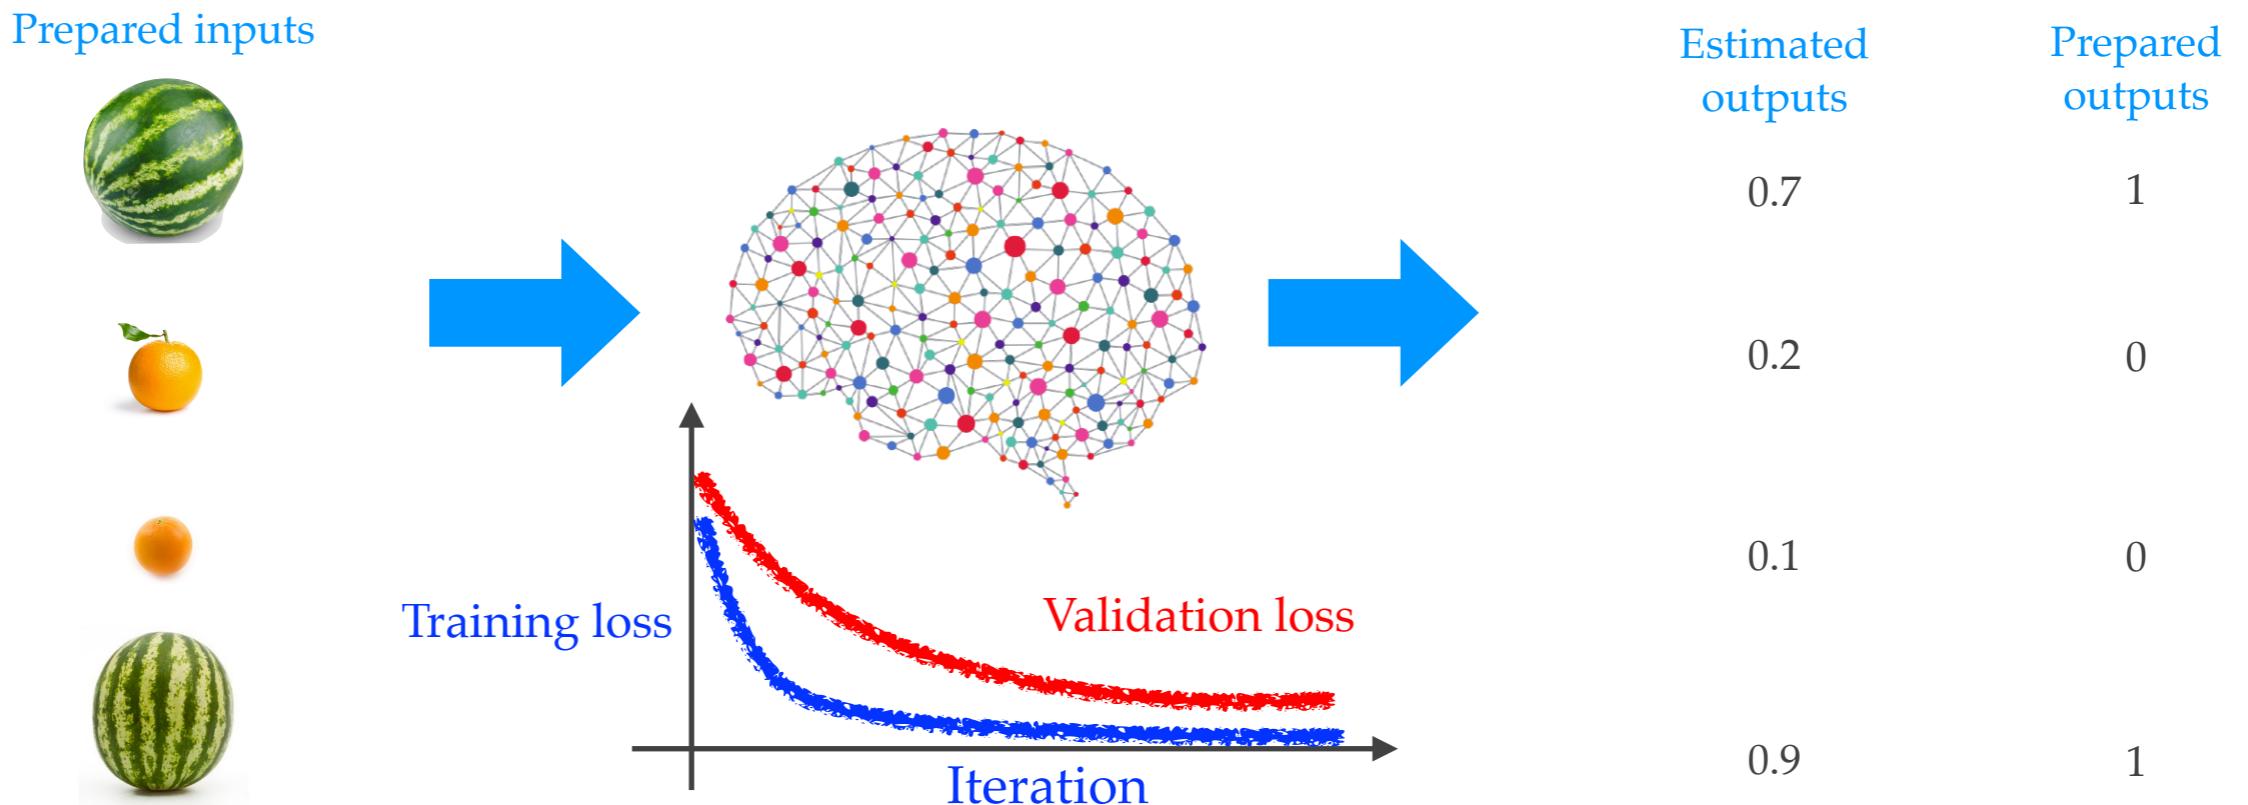


Image Classification Using Deep Learning

- ❖ Training phase: Optimize the weights of a deep neural network



- ❖ Test phase: Perform prediction using the trained neural network

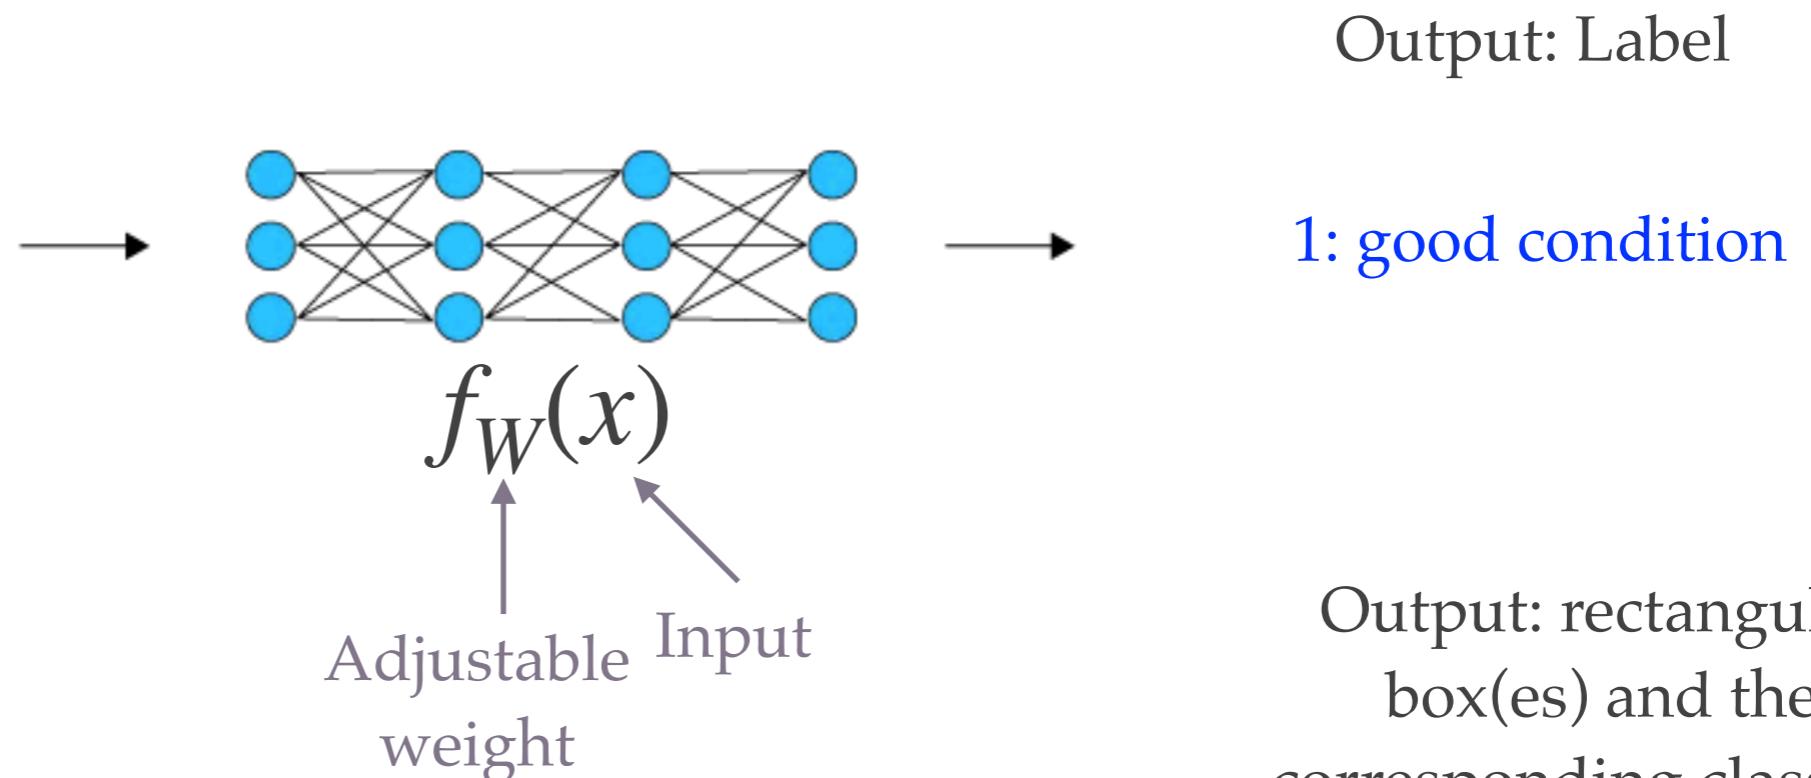


Deep Learning

- ❖ Train a deep neural network (NN) to perform some task from data
- ❖ Think of a deep NN as a universal function approximator

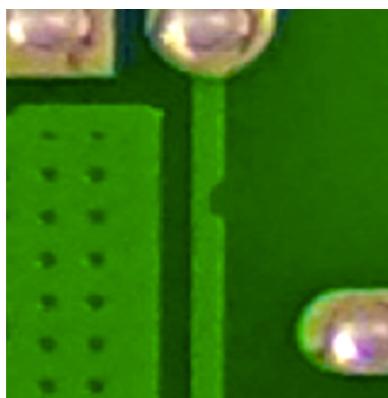
Image classification

Input: image

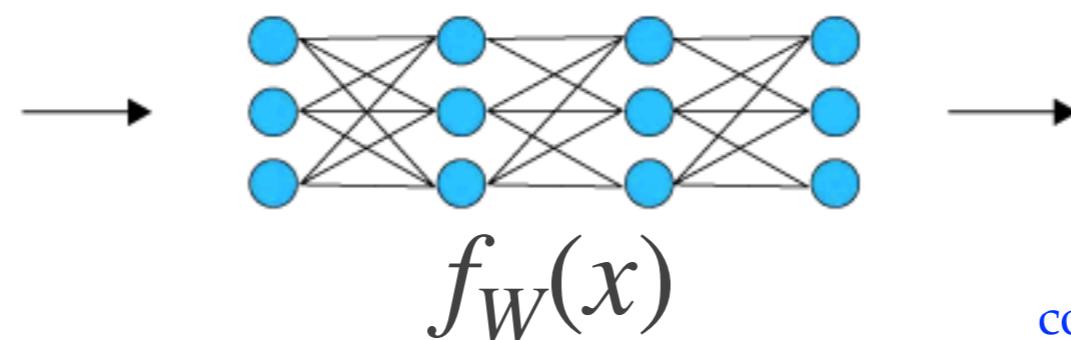


Defect detection

Input: image



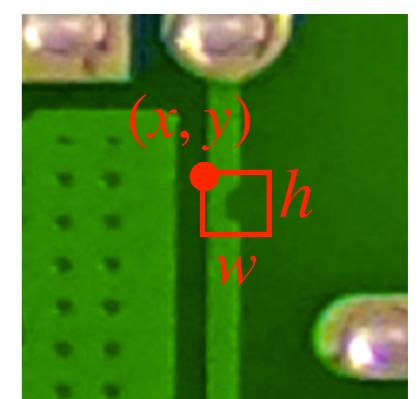
Huang, Weibo, and Peng Wei. "A PCB dataset for defects detection and classification." arXiv preprint arXiv:1901.08204 (2019).



Output: rectangular box(es) and the corresponding class(es)

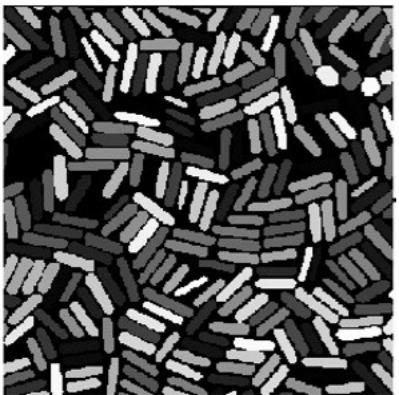
$$\begin{bmatrix} x \\ y \\ w \\ h \end{bmatrix}$$

corresponding class

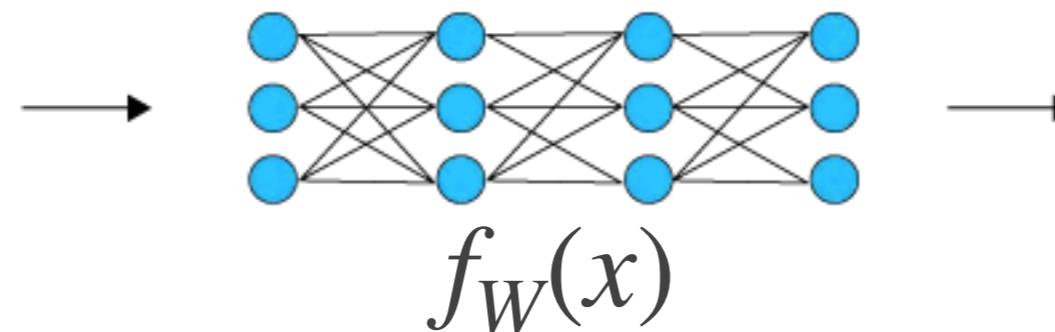


Electron Microscopy Particle Segmentation

Input
EM image



Yildirim, Batuhan, and Jacqueline M. Cole. "Bayesian particle instance segmentation for electron microscopy image quantification." Journal of Chemical Information and Modeling 61.3 (2021): 1136-1149.

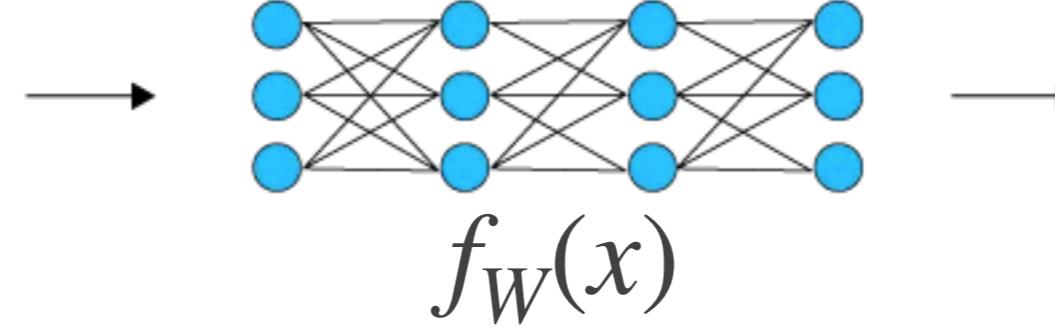
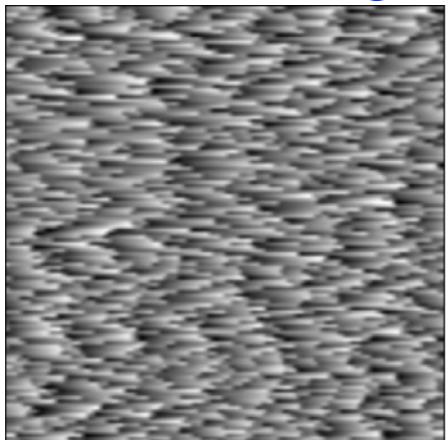


Output
segmentation map

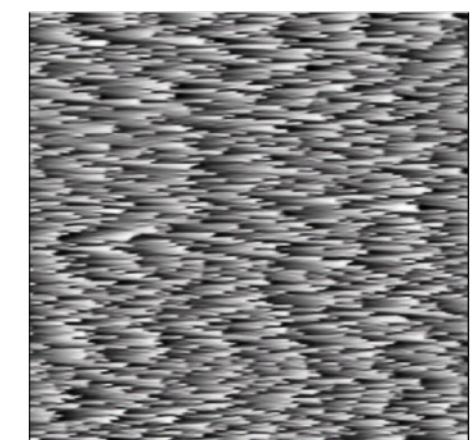


Super-resolution

Input
low-res image

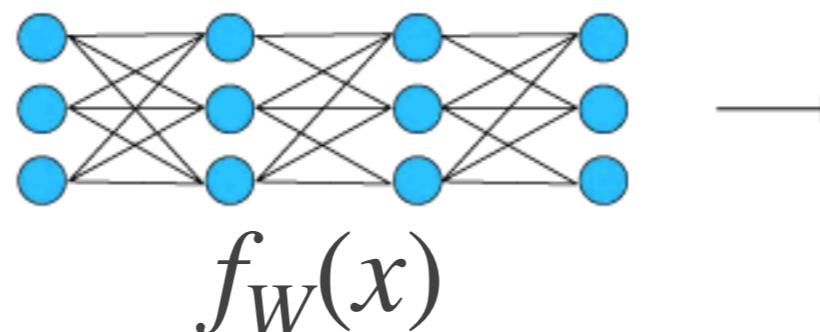
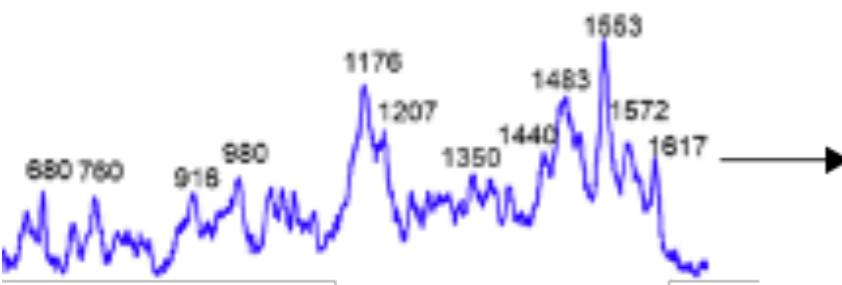


Output
high-res image



Spectrum classification

Input



Output

Category

Synthesis parameter optimization

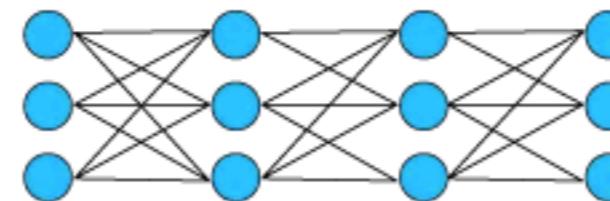
Min, Kyoungmin, et al. "Machine learning assisted optimization of electrochemical properties for Ni-rich cathode materials." *Scientific reports* 8.1 (2018): 1-7.

Input

Experimental parameters

- Composition
- Temperature
- Dopant
- Washing

- Coating
- ICP
- XRD



$$f_W(x)$$



Electrochemical properties

- The 1st discharge capacity
- Capacity retention rate
- Amount of residual Li after initial synthesis

Output

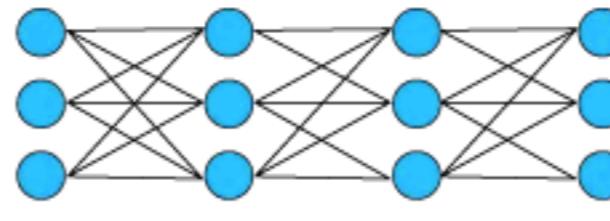
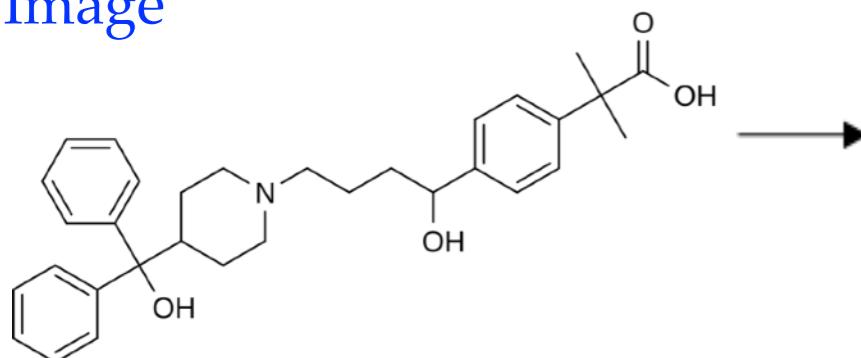
Molecule classification

Input

String of characters



Image



$$f_W(x)$$



2nd-generation
antihistamine
drug

Graph

